

Tejas Gaikwad | Software Engineer

(716) 910-8468 | tejassgaikwad.dev@gmail.com | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

SUMMARY

Software Engineer with 3+ years of experience delivering scalable backend platforms, microservices, and cloud infrastructure for enterprise applications. Expertise spans Java, Python, Spring Boot, FastAPI, AWS, Kubernetes, and modern software architecture, with recent experience building production AI systems powered by LLMs and Retrieval-Augmented Generation (RAG). Recognized for improving application performance, automating deployments, and designing resilient systems that enable rapid product delivery and long-term scalability.

PROFESSIONAL SKILLS

Programming Languages: Java, Python, JavaScript, SQL, Bash

Backend Development: Spring Boot, FastAPI, Node.js, REST APIs, Microservices, API Development, API Integration, Backend Architecture

AI & Machine Learning: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Generative AI, AI Agents, Prompt Engineering, Model Serving, vLLM

Cloud & Infrastructure: AWS (EKS, EC2, IAM, DynamoDB, CloudWatch), Kubernetes, Docker, GitHub Actions, Jenkins, CI/CD Pipelines

Databases: PostgreSQL, MongoDB, Amazon DynamoDB, Redis, SQL Query Optimization, Database Indexing, Data Migration

Software Engineering: Distributed Systems, Scalable Backend Systems, System Design, Performance Optimization, Debugging, Code Reviews, Agile/Scrum

Web & Integration: Chrome Extension Development, OAuth, Third-Party API Integration, Image Processing, Workflow Automation

EXPERIENCE

Software Engineer | DoorDash | NY, USA

Mar 2026 - Present

- Architected a cloud platform on AWS and Kubernetes (EKS) to support production LLM and Vision-Language Model workloads, reducing inference latency through GPU resource optimization.
- Built reusable Terraform modules and GitHub Actions CI/CD pipelines to automate infrastructure provisioning across multi-account AWS environments, reducing deployment time by 30%.
- Optimized model serving with PyTorch, vLLM, and Triton Inference Server by tuning GPU scheduling and autoscaling policies, improving GPU utilization by 30% while lowering compute costs by 18%.
- Developed observability dashboards with Datadog, Prometheus, and AWS CloudWatch to monitor GPU performance, token usage, and service health, reducing troubleshooting time by 35%.
- Designed centralized LLM Gateway, Agent Gateway, and policy-based guardrails using Open Policy Agent (OPA), enabling engineering teams to integrate Generative AI services with consistent security and governance.
- Collaborated with machine learning, platform, and product engineers to deploy AI services using AWS IAM, Kubernetes, and Infrastructure as Code, maintaining 99.9% service availability while supporting multiple production releases.

Software Engineer | Capital One | USA

Jul 2025 - Dec 2025

- Developed Python-based backend services and REST APIs using FastAPI, integrating PostgreSQL and Amazon DynamoDB to support internal applications, contributing to 20% faster API response times.
- Assisted in building Retrieval-Augmented Generation (RAG) pipelines using Python and implementing backend business logic, improving AI-powered search relevance across 10K+ enterprise knowledge records.
- Optimized SQL queries and integrated Redis caching to improve application performance, helping reduce response latency by approximately 25% for frequently accessed endpoints.
- Supported application deployment and CI/CD workflows using Docker, AWS, and Jenkins, contributing to 15% faster release cycles and more reliable deployments.
- Collaborated with a 10+ member Agile engineering team on code reviews, debugging, and third-party API integrations, contributing to the successful delivery of 5+ sprint releases.

Software Engineer | Hexaware Technologies | India

Feb 2022 - Jul 2024

- Developed REST APIs and backend features using Java, Spring Boot, and SQL, improving application response times by 30% while delivering reliable functionality for client-facing applications.
- Refactored key modules from a monolithic application into microservices, simplifying maintenance, improving scalability, & reducing production issues.
- Optimized SQL queries, database indexes, and data migration scripts, cutting report processing time by 20% and improving performance for high-volume business transactions.
- Implemented CI/CD pipelines and automated deployment processes within Agile sprints, reducing release time by 35% and enabling faster delivery of new features.
- Collaborated with a cross-functional team of 10 developers, QA engineers, designers, and project managers to deliver 20+ sprint releases on schedule while quickly resolving technical roadblocks.
- Reviewed code, resolved production issues, and shared best practices with teammates, reducing post-release defects and improving overall code quality across the project.

PROJECTS

Jobleap – [AI-Powered Job Search & Chrome Extension](#) | Node.js, MongoDB, Docker, OAuth, LLM APIs

- Built an automated job application platform connecting four external APIs, authentication, backend services, and a Chrome extension into a unified workflow.
- Automated ATS form completion using LLM-based resume parsing and intelligent field mapping, achieving over 90% submission accuracy through validation.

[Virtual Clothes Try-On](#) | GenAI, IDM-VTON, Node.js

- Developed a Chrome extension that lets shoppers virtually try on clothing directly on retail websites, attracting more than 1,000 early users.
- Integrated server-side image processing and garment rendering with a responsive interface for faster virtual try-ons and reusable uploaded images.

LLM Reasoning Router – [Smart Model Routing System](#) | Python, FastAPI, PostgreSQL, Gemini API, Docker

- Created a routing engine that analyzes request complexity and selects the most suitable LLM using configurable decision rules.
- Added real-time streaming, response quality evaluation, and usage tracking, reducing API costs by approximately 40% while maintaining output quality.

EDUCATION

Master of Science in Computer Science & Engineering | State University of New York Buffalo, USA (Grade: 3.6)

Dec 2025